

# Linux Raw Sockets

And some kernel stuff too...

Jaime Geiger



# Sockets for scrubs (a.k.a. normal people)

*//I can sock it*

*import socket*

*s = socket.socket(socket.AF\_INET,  
socket.SOCK\_STREAM)*

*s.bind(("", 6969))*

*s.listen(123)*

*s.accept()*

*read... write... etc... then close*

# Sockets for programmers (a.k.a. people that are unhealthily obsessed with C)

*//how to socket*

*struct sockaddr\_in serv, cli;*

*int fd = socket(AF\_INET, SOCK\_STREAM, 0);*

*memset(&serv, 0, sizeof(struct sockaddr\_in));*

*serv.sin\_port = htons(6969);*

*serv.sin\_family = AF\_INET;*

*serv.sin\_addr.s\_addr = INADDR\_ANY;*

*bind(fd, (struct sockaddr\*)&serv, sizeof(serv));*

*listen(fd, 5);*

*size\_t len;*

*int clifd = accept(fd, (struct sockaddr\*)&cli, &len);*

*read... write... etc etc. then close*

# As if that wasn't annoying enough...

- We do SOCK\_RAW
- Why bother?
  - Learning and stuff
  - Sniffing is cool
  - Getting closer to the bits

# Socket Buffers: Kernel Magic (50+ members)

```
struct sk_buff {  
    struct sk_buff *next, *prev; ...  
    struct sock *sk;  
    struct net_device *dev; ...  
    unsigned int len, data_len;  
    __u16 mac_len, hdr_len; ...  
    union {__wsum csum; struct {__u16 csum_start; __u16  
        csum_offset;};}; ...  
    __be16 protocol;  
    __u16 transport_header, network_header, mac_header; ...  
    sk_buff_data_t tail, end;  
    unsigned char *head, *data;  
}
```

# How does this all happen?

- Network card gets packet, sends interrupt to the kernel
- Device driver → `netif_rx(skb)` →
- `enqueue_to_backlog(skb, cpu, &qtail)` →
  - if queue isn't full and flow rate is not met then add to per-cpu processing queue
  - calls soft irq (`NET_RX_SOFTIRQ`)
- `dev_add_pack` → `ptype_head(pack_type)`
  - Assigns packet type structure to the skb queue item
- `net_rx_action(softirq_action_ptr)` ...

# net\_rx\_action(softirq\_action\_ptr)

- Confusing interrupt handler
- struct napi\_struct \*n;
  - Will contain the queued skb
- n = list\_first\_entry(&sd->poll\_list, struct napi\_struct, poll\_list);
  - Gets queued packet
- Runs handler functions for packet types
  - ptype\_all (ETH\_P\_ALL)
  - ptype\_base[<type hash>] (ETH\_P\_\*)
  - Each type is passed further up their respective protocol stacks

# Sending Raw Packets

- You have to manually fill out EVERY FIELD
- There is some header help
  - Ethhdr, iphdr, udphdr, and tcphdr structs, combine them to make entire frames:

```
struct __attribute__((__packed__)) udpframe {  
    struct ethhdr ehdr;  
    struct iphdr ip;  
    struct udphdr udp;  
    unsigned char data[ETH_DATA_LEN - sizeof(struct udphdr) - sizeof(struct iphdr)];  
};
```

\*\*yeah where is that FCS, don't know, but things packed into this correctly will send :)



# Sending Raw Packets

- sendto requires a sockaddr\_ll (link layer) to send raw frames
  - Source MAC address, protocol family (PF\_PACKET), hwaddr len, and interface index
  - ioctl(sockfd, SIOCGIFINDEX, sifreq) gets us the interface index in sifreq->ifindex
- Checksums suck, UDP doesn't require one, just set it to 0x0000
  - UDP is just easier to send in general, no handshakes to work out
- If you send raw UDP responses to sniffed frames you won't be able to receive them unless you block icmp port unreachable messages

# Packet Sniffing

- Send an ioctl to the interface to set flags to allow promiscuous mode:

```
struct ifreq *sifreq;
```

```
int sockfd = socket(PF_PACKET, SOCK_RAW,  
htons(ETH_P_IP));
```

```
strncpy(sifreq->ifr_name, "eth0", IFNAMSIZ);
```

```
ioctl(sockfd, SIOCGIFFLAGS, sifreq);
```

```
sifreq->ifr_flags |= IFF_PROMISC;
```

```
ioctl(sockfd, SIOCSIFFLAGS, sifreq);
```

- Grab packets:

```
recvfrom(sockfd, ...)
```

# PF\_PACKET?

- Gets the specified packet type (ETH\_P\_IP) directly off of the interface driver instead of letting it go farther up the protocol stack
- Why is this good?
  - The sk\_buffs get copied, so you can still browse the internet while you sniff! (literally in skb\_copy)
  - The packet has not been modified or filtered in any way higher levels of the networking stack

# Berkley Packet Filter (BPF)

- So say you don't want all of the IP packets...
  - (Trust me, you don't)
- Sockets can have filters applied to them with `setsockopt`
  - Filters at the socket level! No need to filter through in *your* code, **that's really slow**
- Use `tcpdump` to generate a filter

# Making and Applying socket\_filters

```
$ sudo tcpdump -dd udp and port 31337
```

```
{ 0x28, 0, 0, 0x0000000c },
{ 0x15, 0, 6, 0x000086dd },
{ 0x30, 0, 0, 0x00000014 },
{ 0x15, 0, 15, 0x00000011 },
{ 0x28, 0, 0, 0x00000036 },
{ 0x15, 12, 0, 0x00007a69 },
{ 0x28, 0, 0, 0x00000038 },
{ 0x15, 10, 11, 0x00007a69 },
{ 0x15, 0, 10, 0x00000800 },
{ 0x30, 0, 0, 0x00000017 },
{ 0x15, 0, 8, 0x00000011 },
{ 0x28, 0, 0, 0x00000014 },
{ 0x45, 6, 0, 0x00001fff },
{ 0xb1, 0, 0, 0x0000000e },
{ 0x48, 0, 0, 0x0000000e },
{ 0x15, 2, 0, 0x00007a69 },
{ 0x48, 0, 0, 0x00000010 },
{ 0x15, 0, 1, 0x00007a69 },
{ 0x6, 0, 0, 0x0000ffff },
{ 0x6, 0, 0, 0x00000000 },
```

```
struct sock_fprog filter;
filter.len = 20;
filter.filter = bpf_code;
setsockopt(sockfd, SOL_SOCKET,
SO_ATTACH_FILTER, &filter, sizeof(filter));
```

- The above will set the bpf filter on the socket.
- Stored in sock->sk\_filter
- Now recvfrom(sockfd, ...) will only get IP packets destined for UDP port 31337

```
struct sock_filter bpf_code[] = {
    { the, filter, from, above },
};
```

# Watershell

- A practical C implementation of most of this
- Run commands through iptables!
  - Sniffing pulls packets directly off the interface before iptables has a chance to examine them
  - Packets are passed through iptables after they are picked up by the interface
  - Special port and keyword used to trigger command execution via network (just a system() call for now)
- <https://github.com/jgeigerm/watershell.git>

# You can do all of this with python

- Yeah python makes all of this a lot easier, but C is more hardcore or whatever
- Plus I like compiled binaries I can drop onto systems
  - You can do that with python too
    - ...